

SECURITY IN A DISTRIBUTED APPLICATION

Brian Cheung and Sam Wang

Electrical and Computer Engineering, Cockrell, University of Texas at Austin

Abstract

Container orchestration technologies like Kubernetes provide a simple way to deploy, scale, and manage containerized applications [2]. However, as with any production-level system, these distributed systems must be properly secured at each appropriate level. In this project, we examine common vulnerabilities of distributed applications and demonstrate how we can help mitigate them through Pod Security Policies and network proxies like Envoy. We first implement Pod Security Policies to improve the base security of the application itself by limiting the pod's capabilities and only allowing appropriate privileges. Then we implement Envoy with Open Policy Agent to reinforce the network security by only accepting specific requests based on client privileges. As shown through our examples, these additional layers of security can help mitigate common vulnerabilities.

Motivation

With the increasing popularity of distributed applications deployed with container orchestration technologies like Kubernetes, developers must understand how to reinforce the security of these distributed applications. Additionally, we must consider additional layers of security, which depend on the type of application and its exposure to other services. Often times, services communicate with other services through a networks, so a network proxy that filters incoming and outgoing requests can help protect the service itself along with the clients who use it. Securing an application at multiple levels allows developers to reduce the overall attack surface by making it more difficult for attackers to find vulnerabilities.

Common Vulnerabilities and Threats

Data Breach/Loss: Attackers will target sensitive/confidential/personal data that some application holds.

Denial of Service: This type of attack can appear in the form of a Denial of Service (DoS) or Distributed Denial of Service (DDoS). In a DoS attack, one machine will start the attack and prevents legitimate users from using some service/application. A DDoS attack is when an attacker will take control of multiple systems to use to take down some application.

Vulnerable Cloud API and Systems: By accidentally exposing an API to the public, attackers can exploit that API and can attack a system.

Malicious Insider(s): When someone attacks from within the application, like an employee, system administrator, business partner, etc. This type of attack can get access to sensitive information and can spread malware throughout the entire system or even to clients.

Account Hijacking: When an account's credentials are compromised by a hacker and used for purposes against the owner's will.

Lack of Due Diligence: Due diligence is the practice of knowing and applying the best security practices. The lack of a good foundation and implementation of best security practices leaves a service extremely vulnerable to attacks. [3]

Pod Security Policies (PSP)

Pod Security Policies (PSP) are policies that define rules/conditions that a pod must run with in order to function within the cluster of pods. This ensures that pods can and will run only if it maintains the correct privileges and resources defined in the PSPs. When a pod is deployed, the PSP acts as a gatekeeper that will compare the pod security configuration to what is defined in the PSP. Some examples of how PSPs can limit pod behaviors include, preventing privilege escalation, restricting pods from accessing host namespaces, filesystem, and networks, restrict the amount of user/groups that a pod can run, and more.

Examples

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod/non-root-debian-5b8bdd6c5-cjt14	2/2	Running	0	2m18s	10.1.56.103	bcheung	<none>	<none>
pod/root-debian-69dcfd6f4c-b8kgz	0/1	CreateContainerConfigError	0	2m17s	10.1.56.104	bcheung	<none>	<none>
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR		
service/kubernetes	ClusterIP	10.152.183.1	<none>	443/TCP	40h	<none>		
service/non-root-debian-svc	NodePort	10.152.183.21	<none>	80:30080/TCP	2m17s	appnon-root-debian		
service/root-debian-svc	NodePort	10.152.183.31	<none>	80:30001/TCP	2m18s	approot-debian		
NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR	
deployment.apps/non-root-debian	1/1	1	1	2m18s	non-root-debian	localhost:32000/app_non_root_debian:k8s	appnon-root-debian	
deployment.apps/root-debian	0/1	1	0	2m17s	root-debian	localhost:32000/app_root_debian:k8s	approot-debian	
NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR	
replicaset.apps/non-root-debian-5b8bdd6c5	1	1	1	2m18s	non-root-debian	localhost:32000/app_non_root_debian:k8s	appnon-root-debian,pod-template-hash=5b8bdd6c5	
replicaset.apps/root-debian-69dcfd6f4c	1	1	0	2m17s	root-debian	localhost:32000/app_root_debian:k8s	approot-debian,pod-template-hash=69dcfd6f4c	

Fig. 1: Current Resources with restrictive PSP applied.

Conditions:					
Type	Status				
Initialized	True				
Ready	False				
ContainersReady	False				
PodsScheduled	True				
Volumes:					
default-token-6clvb:					
Type:	Secret (a volume populated by a Secret)				
SecretName:	default-token-6clvb				
Optional:	false				
QoS Class:	BestEffort				
Node-Selectors:	<none>				
Tolerations:	node.kubernetes.io/not-ready:NoExecute for 300s				
	node.kubernetes.io/unreachable:NoExecute for 300s				
Events:					
Type	Reason	Age	From	Message	
Normal	Scheduled	<unknown>	default-scheduler	Successfully assigned default/root-debian-69dcfd6f4c-b8kgz to bcheung	
Normal	Pulled	7s (x2 over 7s)	kubelet, bcheung	Container image "localhost:32000/app_root_debian:k8s" already present on machine	
Warning	Failed	7s (x2 over 7s)	kubelet, bcheung	Error: container has runAsNonRoot and image will run as root	

Fig. 2: Describing Pod with root as default user.

```
bcheung@bcheung:~/Documents/EE379K-Final-Project$ microk8s.kubectl exec -it non-root-debian-5b8bdd6c5-cjt14 bash
appuser@non-root-debian-5b8bdd6c5-cjt14:/ $ ls
bin boot dev etc home lib lib64 media mnt opt proc root run sbin srv sys tmp usr var
appuser@non-root-debian-5b8bdd6c5-cjt14:/ $ cd root
bash: cd: root: Permission denied
appuser@non-root-debian-5b8bdd6c5-cjt14:/ $ su root
Password:
su: Authentication failure
```

Fig. 3: Attempting Privilege Escalation with PSP applied.

```
appuser@bcheung:~/Documents$ curl -i -H "Authorization: Bearer 'SSMIA_TOKEN'" http://SERVICE_URL/people
HTTP/1.1 200 OK
content-type: application/json
date: Sat, 14 Dec 2019 19:51:57 GMT
content-length: 151
x-envoy-upstream-service-time: 0
server: envoy

[[{"id":"1","firstname":"John","lastname":"Doe"}, {"id":"2","firstname":"Jane","lastname":"Doe"}, {"id":"498881","firstname":"Charlie","lastname":"Doe"}]]
appuser@bcheung:~/Documents$ curl -i -H "Authorization: Bearer 'SSMIA_TOKEN'" -d '{"firstname":"Charlie","lastname":"Doe"}' -H "Content-Type: application/json" -X POST http://SERVICE_URL/people
HTTP/1.1 403 Forbidden
date: Sat, 14 Dec 2019 19:52:24 GMT
content-length: 0
server: envoy
content-length: 0
```

Fig. 4: Envoy with OPA Policy Guest.

```
appuser@bcheung:~/Documents$ curl -i -H "Authorization: Bearer 'SSMIA_TOKEN'" http://SERVICE_URL/people
HTTP/1.1 200 OK
content-type: application/json
date: Sat, 14 Dec 2019 19:52:33 GMT
content-length: 151
x-envoy-upstream-service-time: 0
server: envoy

[[{"id":"1","firstname":"John","lastname":"Doe"}, {"id":"2","firstname":"Jane","lastname":"Doe"}, {"id":"498881","firstname":"Charlie","lastname":"Doe"}]]
appuser@bcheung:~/Documents$ curl -i -H "Authorization: Bearer 'SSMIA_TOKEN'" -d '{"firstname":"Charlie","lastname":"Doe"}' -H "Content-Type: application/json" -X POST http://SERVICE_URL/people
HTTP/1.1 200 OK
content-type: application/json
date: Sat, 14 Dec 2019 19:52:41 GMT
content-length: 55
x-envoy-upstream-service-time: 0
server: envoy

{"id":"727887","firstname":"Charlie","lastname":"Doe"}
```

Fig. 5: Envoy with OPA Policy Admin.

Envoy

Envoy is an open source network proxy that can be customized to support TCP/HTTP proxy, TLS authentication, load balancing, etc. Envoy aims to achieve network and application transparency and complete communication between the two. In our report, we demonstrated Envoy's ability to secure network traffic with HTTPS and SSL certificates as well as Envoy's external authorization filter functionality, specifically applying a policy written with Open Policy Agent [1].

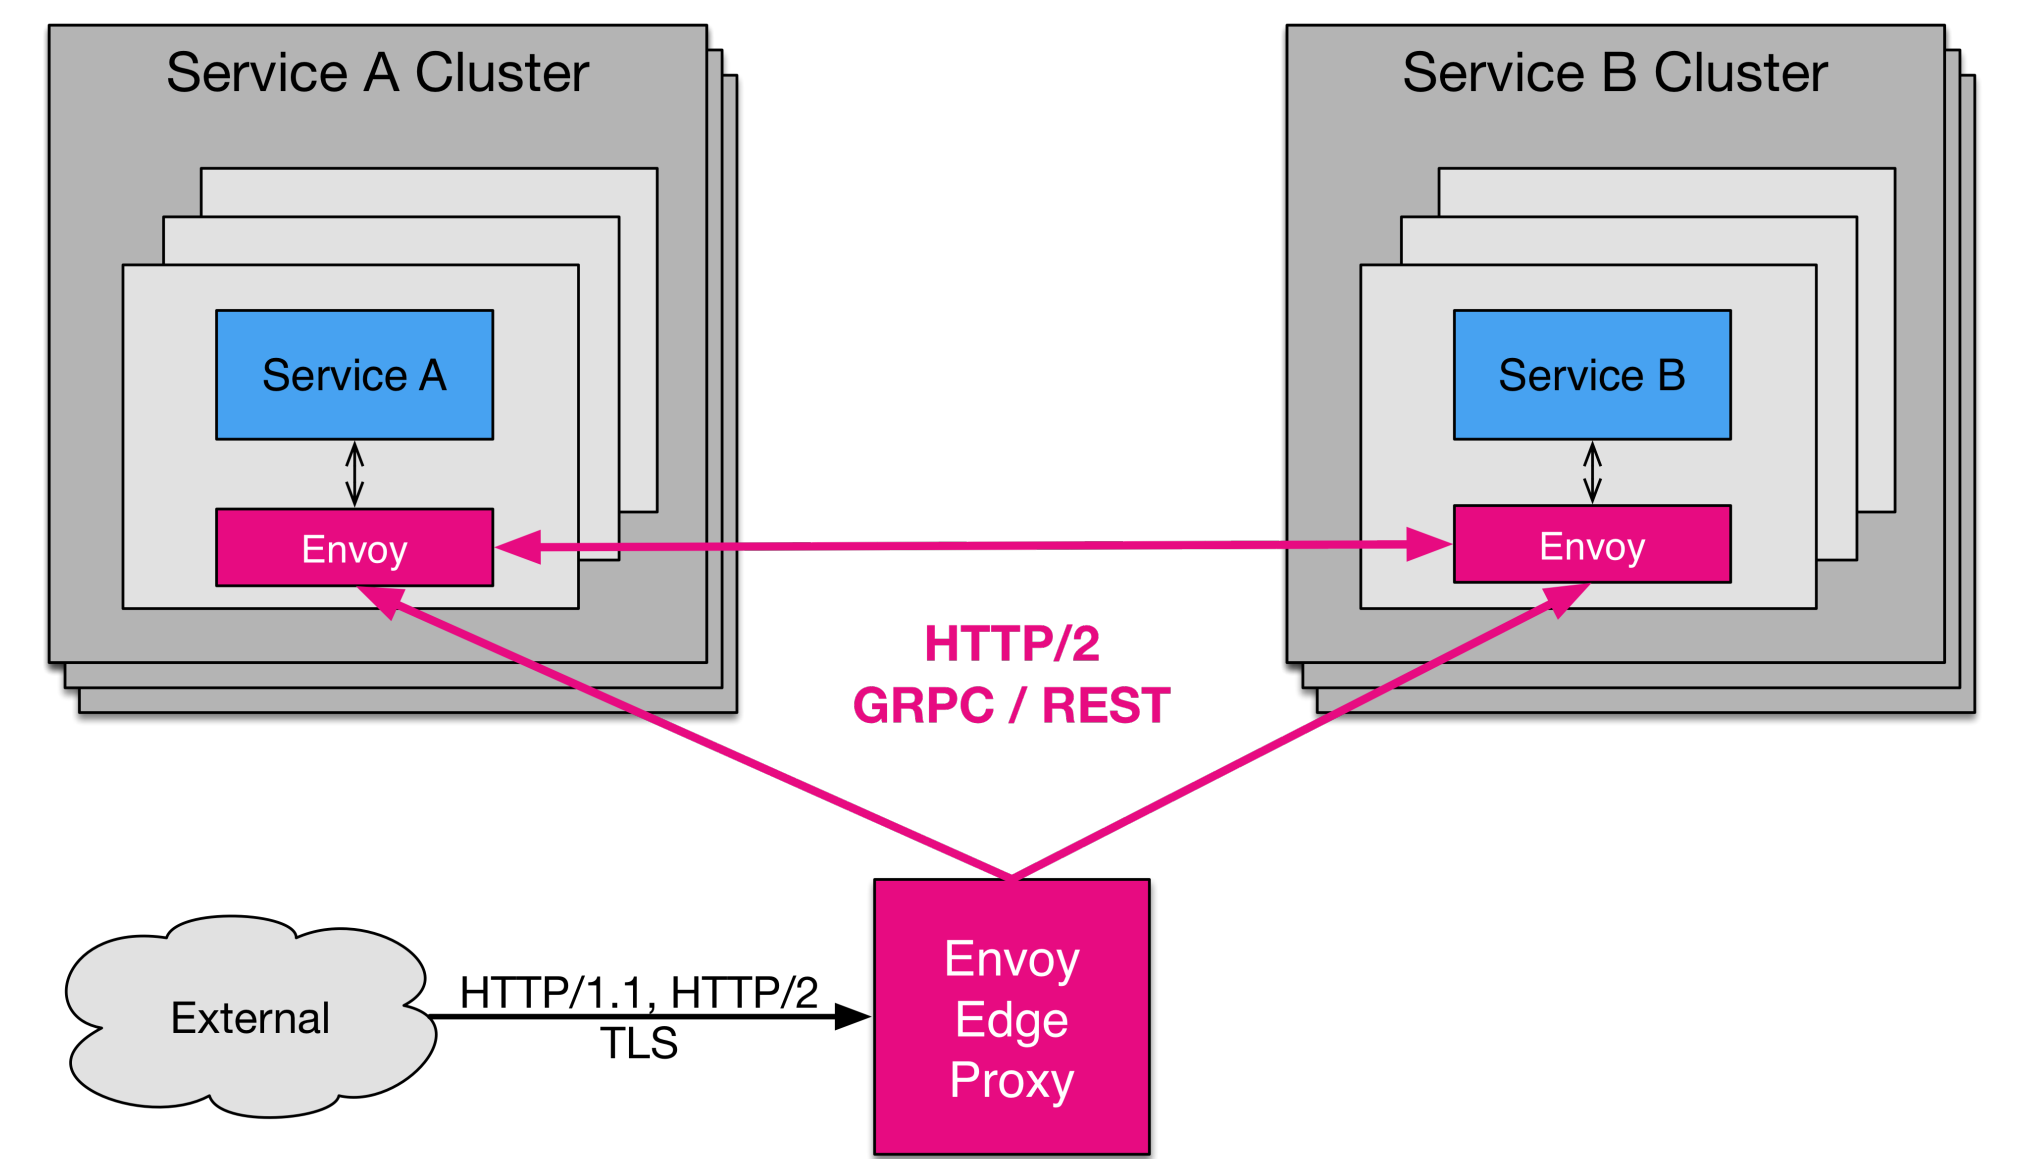


Fig. 6: Envoy Diagram.

Comparison

While Pod Security Policies and network proxies have similar goals of maintaining/strengthening the security of an application, they do have different use cases and have unique strengths when it comes to application/network security. Pod Security Policies secure the base application and the Kubernetes cluster that the application is hosted on. On the other hand, a network proxy is the middleman between the client and the network and thus is met with every network request. Furthermore, designing a network proxy that can effectively handle these network requests, whether that be filtering out malicious requests, or assigning security policies onto API requests is extremely important in protecting the overall security integrity of the application as well as reducing the load that the rest of the application has to deal with. Network proxies can be viewed as the outer line of defense, while Pod Security Policies act as the base-line defense that protects the system if a threat manages to bypass the network proxy. Both services add a layer of security that helps mitigate threats at different levels of the application, which ultimately minimizes the attack surface of the application.

References

- [1] Envoy. "What is Envoy". https://www.envoyproxy.io/docs/envoy/latest/intro/what_is_envoy.
- [2] Kubernetes. "Production-Grade Container Orchestration". <https://kubernetes.io/>.
- [3] Pericherla Satya Suryateja. *Threats and Vulnerabilities of Cloud Computing: A Review*. https://www.researchgate.net/publication/324562008_Threats_and_Vulnerabilities_of_Cloud_Comput